

The decision system that learns your judgement.

Citra recommends the call on every case — then learns from every accept, reject and the **why** your officers record. Delivered as a Decision App, a Decision API and an embeddable recommendation UI — all authored in plain English. **People leave, models change — the judgement stays.**

THE PROBLEM — YOUR COMPANY FORGETS EVERY DECISION IT MAKES

A senior officer rejects a ₹35 lakh loan — the GST filings don't match the bank statements. Her reasoning is sound, and it lives only in her head. Two years on, a near-identical file lands at another branch. **She's gone. It's approved.** Underwriting is genuinely hard: statements, GST, ITR, borrower history, bureau data, collateral, and the tail of things a good officer notices. All that diligence happens, a call is made — and the institution keeps the verdict while losing the reasoning that produced it.

The people door

When a senior officer leaves, the judgement leaves with them. Onboarding restarts the clock, and consistency between branches drifts.

The model door

Model weights hold patterns, not reasons. No model can cite the past record that justified a decision — and none survive a model swap with their reasoning intact.

THE OUTCOME

3 fabricated files caught before payout. ₹15 lakh that never leaves the bank — per 100 DSA-sourced applications.

MODELLED on a 100-file cohort. The control pays for itself above 1 file in 1,100 — the experiment behind it is on page 2.

HOW IT WORKS — THE LEARNING LOOP

01

Author in plain English

Your credit or ops lead describes the workflow — "loan underwriting", "recovery", "claims review". Citra generates the Decision App and its API in days.



02

Citra recommends

Each case arrives assembled and scored, with a recommended action and the reason for it — grounded in your own records and your own past decisions.



03

Your officer decides

Approve or override, and record the why. The AI never acts on its own — auto-commit is opt-in, later, only where you choose.



04

The judgement compounds

Every why is written to a ledger you own. Retrieval grounds the next decision, and every proven call becomes a free, verified training label.

CITRA MEMORY — the decision ledger you own

case · evidence · recommendation · override · the WHY · outcome — the golden source that outlives any model

HOW IT REACHES YOUR PEOPLE — THREE SURFACES, ONE INTELLIGENCE

Decision App

The ready-made workspace — case-working pages, live dashboards and a plain-English copilot. Live on day one, nothing to build.

Decision API

Every recommendation, score, reason and learning loop served over REST, with SDKs. Call it from any system you already run.

Embeddable recommendation UI

Drop the recommendation and its reasoning straight into your existing LOS, CRM or core screens — web, desktop, Android, iOS. Your officers never change tools.

WHAT IT DELIVERS — THREE SAVINGS THAT COMPOUND, IN THE ORDER THEY ARRIVE

NPAs cut

Losses prevented — this is the one that pays

Higher-quality recommendations catch what slips through: loans heading for NPA, and fraudulent claims stopped before the money leaves. **One NPA or fraudulent claim stopped can outweigh the annual licence on its own** — and it compounds, because the more judgement it captures the more it catches.

40–70%

Effort removed — capacity released, not headcount

Low-complexity cases resolve in one click; complex ones keep human approval on an AI recommendation. Your officers stop assembling context and spend the time judging the files that need judgement. On a ~2M-decision workflow that is **₹4–10 Cr of capacity a year** — most partners spend it absorbing growth, not cutting heads.

Days

To first app — not a 3–6 month build

Each Decision App and API ships in days. The IT or SI integration project you'd otherwise fund is saved on day one.

You own it

The ledger, and everything in it

Every decision on record with its reasoning — a governance trail for auditors and regulators, and an asset that survives every model change.

Losses prevented come first, then operations capacity released — build-time avoided is the day-one rupee on top. A 100-person ops team (~₹15 Cr/yr loaded) working 40–70% less manual effort releases ₹6–10 Cr of capacity against a ₹50 lakh–₹1 Cr licence plus metered usage — capacity most lenders put into higher volume and shorter turnaround rather than into cuts. **Together they cover the licence many times over, and payback lands within year one.**

Safe and sovereign by design

- Citra recommends; your expert approves. A wrong suggestion is just a rejected recommendation — and the rejection teaches it.
- A governed ontology bounds what the AI may build and do; every write is schema-validated before it runs.
- **Two deployment options. Citra Private SaaS** — a dedicated single-tenant instance we run for you. **BYOC** — deployed inside your own private cloud, you hold the keys and the infrastructure. Zero data egress either way; open-source models metered in-tenant.
- Decides across SAP, SQL Server, Salesforce and your homegrown tail — no rip-and-replace, no vendor lock.
- Loss prevention is built in: reused documents and tampered statements are flagged automatically.

The 90-day design-partner pilot — bounded, paid, recommend-only

Scope: one workflow you choose — underwriting, recovery, claims or invoice review.

Success criteria, agreed in writing before we start: decision-quality lift, effort reduced, time-to-first-app, zero incorrect writes.

Safety: recommend-only — nothing writes to your systems without your officer's approval. Autonomy is opt-in, later, only if you ask.

Paid: a real but reduced pilot fee. Paying partners give honest feedback; free pilots get abandoned.

Founder-led: hands-on in your environment weekly, fixes in days.

Then: annual licence from ₹50 lakh ACV, banded by scale. Private cloud, not per-seat.

What that means concretely. We carry the delivery risk — criteria fixed before we start, founder-led execution, no autonomous writes. You keep the outcome — the working app, the decision ledger and everything in it, and a referenceable result — whether or not you continue with us.

Let's build the memory together. **Start with one workflow.**

Your credit policy is complete. It is also **silent** — and that is where the money leaves.

A worked example from Indian retail underwriting. This page shows where the money sits; the next page shows what we actually measured, and what we did not.

MEASURED observed in a controlled run **MODELLED** arithmetic on stated assumptions, not yet observed

WHAT "DECISION MEMORY" MEANS HERE

When your officer overrides a recommendation, Citra records the decision, the reason and the owner. Corrections that repeat across officers harden into a **learned judgement** — a rule your written policy never contained, attributed to the people who formed it, stored in a ledger you own. Citra then applies it to the next matching case. **A learned judgement is the memory.**

THE ONE WE TESTED — LEARNED JUDGEMENT C-002

"DSA-sourced files get employment verified with the employer directly — the submitted document set is not enough."

Formed from corrections by **three named credit officers**. It appears in **no policy document** — so any behaviour change cannot have come from reading the corpus. Its scope is **one facet — sourcing_channel: dsa**, and only one: a judgement's scope is the intersection of the labels on every correction that formed it, and the channel is the only thing all three teaching cases shared.

The file that passes every gate — a DSA-sourced personal loan

\$4	Bureau band. Below 650 declines unless fully secured	Clears at 726
\$2.2	Income floor. ₹40,000/month net, self-employed	Clears
\$3.1	FOIR cap by income band, 45% to 60%	Clears at 58%
\$5	Documents. ITR two years, bank statement twelve months	Complete
\$1	Scope. Applies to branch, DSA, digital and bancassurance alike	No channel rule

\$5 is a collection requirement, not a corroboration requirement: it says get the ITR. It never says phone the employer.

Why better modelling cannot close this

A DSA is paid on disbursal, not on repayment. Experienced officers know what that incentive produces and apply extra scrutiny. **No credit policy we have seen says so.** Here is everything an underwriting system actually receives per case:

application_id · customer_id · product · amount_requested · tenure_months · sourcing_channel · income_proof_type · itr_declared_income · ltv_percent · foir_percent · status

There is no field for **employment verified**. The closest records only that a document of a given type sits on file — nothing distinguishes **the documents are present** from **the documents are true**. **The signal is not in the columns. It is in a column nobody thought to create.**

Where the money is — the same 100 files, with memory off and with memory on **MODELLED** funnel proportions and catch counts are illustrative arithmetic on a 100-file cohort

100 DSA-sourced applications

run through the policy gates

\$4 bureau · \$2.2 income · \$3.1 FOIR · \$5 documents

BRANCH A — ~70 declined

Dead on a hard rule: bureau 617, FOIR 95%, income below the floor.
A learned judgement cannot change a verdict here — the file fails policy with memory on or memory off.

No money is recoverable in this branch.

BRANCH B — ~30 clear every gate

Recommended for approval — correctly, under the policy as written.
\$5 asked for documents. It never asked for a phone call.
This is where the money is — and our controlled run did not sample it.

MEASURED Where our controlled run lived

14 of the 19 tested cases sat in Branch A.
That is why no final decision changed in the run: with the file already failing a hard rule, a verdict change was structurally impossible. What changed was the reasoning, not the outcome.

B1 · MEMORY OFF

Nothing prompts the extra step.
All 30 disburse. The employer is never contacted.
Loss lands 6–18 months later as principal, provisioning, collection.

B2 · MEMORY ON

C-002 fires and adds one step: verify employment directly.
27 confirmed, cleared in a day.
3 fabricated, caught before payout.
₹15 lakh never leaves the bank.

What the test showed — and what it did not.

We took nineteen DSA-sourced applications and ran each one twice, with everything identical except one thing: whether learned judgement C-002 was **switched on** or **switched off**. We also ran non-DSA files, to check the opposite — that the judgement stays quiet on the cases it was never meant to touch.

What we measured MEASURED

14 vs 1

The judgement was used on 14 cases with memory on. With it off, just 1.

1 in 2,000

The chance of a swing that size happening by luck alone ($p = 0.0005$).

19 of 19

It fired on every DSA file — the cases it was built for.

0 of 2

It stayed silent on the non-DSA files — the ones it was not.

Put simply: turning the memory on changed how the system reasoned, it did so on every case it was meant to, and never on a case it was not.

The null result — the part worth trusting

We seeded four learned judgements. Three merely restated the written SOP. **Retired, they changed nothing** — the system reached the same conclusion without them, because it can simply read the rule. They fired, they were cited, they added no value.

And no approve-or-decline verdict changed in the run. Fourteen of the nineteen cases sat in **Branch A** on the previous page — already failing a hard rule — so no verdict could have changed. What changed was the reasoning: an extra step nobody had written down. **Branch B, where the money is, was not covered by this run at all.**

Any vendor can tell you their system learns. Fewer will tell you which of their own features measurably does nothing.

What it costs to be right MODELLED

How often must the check catch something to pay for itself?

Cost of one check ÷ (loan size × your loss-given-default). A catch on a ₹5 lakh loan at 70% LGD avoids about **₹3.5 lakh**.

Cost assumption	Per file	Break-even
Direct cost of the call only	₹300	0.09%
All-in, loading drop-off from the extra day	₹5,000	1.4%

So the check pays for itself if roughly **1 file in 1,100** is misstated — or 1 in 70 on the harshest cost assumption. Employment misstatement on DSA files runs in the low single-digit percent, so you clear both bars comfortably. **Put your own numbers in: it takes three inputs, and you hold all three.**

Three questions we expect you to ask — and our straight answers

"We already have an LLM with RAG over our policies, and copilots for our officers."

Then for *what does the policy say*, you are done, and we would not try to displace it. Three things that setup structurally cannot do: it **leaves no asset behind** — after two years of copilot usage you own nothing; it has **no outcome loop**, because nobody tells a copilot the loan went NPA fourteen months later; and it produces **no record a regulator accepts**, because a suggestion in a chat window is not a decision with a rule, a rationale and an owner attached.

"We'll train a model on our own decision history instead."

A model trained on the eleven fields shown on the previous page can learn that DSA files in this band go bad more often. Its output is a risk score nudged down a few points — which **prices the uncertainty rather than resolving it**, and leaves the officer holding a file that still looks compliant on every field the system has. C-002 outputs a step: phone the employer. **No model can recommend acquiring data that exists in no column of its training set.**

"If your system finds gaps in our policy, we'll just fix the policy."

You should — and if C-002 holds up on your book, you must. **Promoting a learned judgement into written policy is the intended destination, not a failure of the system.** C-002 reaches the same population a blanket policy line would; what it adds is everything around the rule. **It exists at all** — you can only write the clause once you know it exists. **It is attributed** to three named officers, with the cases one click away. **It is measured from its first outing**, where a policy clause ships unmeasured. And **it is reversible in a click**, where a revision takes another committee.

Proving it on your book — what a 90-day pilot has to return

How many files get diverted. Of the files approved with memory off, how many switch to "verify employment" once it is on. Measured directly.

How much officers disagree. On matched case profiles, how far decisions vary between officers and branches — before and after. Visible in week one.

A pilot exists to convert Branch B from arithmetic into a measured number. It needs one input we cannot supply: your own DSA misstatement rate, which your credit risk team already knows from NPA post-mortems.

How many of those turn out bad. Of the files diverted, how many fail the check. This is the number the whole return rests on — and it becomes yours, not ours.

How fast the first app is live. Built by your own business team in plain English, in days rather than a two-quarter IT build. Provable on day one.

What this note does not claim

This result covers one judgement, one application, nineteen cases, one customer environment. It shows that a learned judgement changes how the system reasons where the written policy is silent, and stays quiet where it is not. It is **not** evidence that decision memory helps in general, and we would not present it as such. The run tested **one facet — the sourcing channel — because that is C-002's entire scope**; it is not evidence about how well multi-facet scopes discriminate, which we have not measured. Every rupee figure here is arithmetic on stated assumptions — the break-even uses one loan size and one flat loss rate, where a real book has a spread of both. A credit head who cannot find the limits of a vendor claim in the vendor's own document assumes they were hidden.